

Week 6: File Ingestion and Schema Validation

1. Objective

The goal of this assignment was to:

- Ingest a large CSV dataset (credit card transactions) using multiple frameworks: **Pandas**, **Dask**, **Modin**.
 - Clean column names and validate them against a predefined **YAML schema**.
 - Write the cleaned data to **pipe-separated compressed (.gz) files**.
 - Compare computational efficiency and output file sizes.
-

2. Dataset

- **File:** creditcard.csv
 - **Rows:** 284,807
 - **Columns:** 31
-

3. Methodology

1. **Reading the file**
 - **Pandas:** standard single-threaded read.
 - **Dask:** parallel read with assume_missing=True to handle dtype inference.
 - **Modin (Ray backend):** parallelized Pandas.
 2. **Column cleaning**
 - Removed whitespace and special characters.
 - Replaced spaces with underscores.
 3. **Schema validation**
 - Created a **YAML file** containing column names and CSV separator.
 - Validated ingested data for both **row count** and **column names**.
 4. **File writing**
 - Wrote **pipe-separated (|) .gz compressed files** for all three frameworks.
-

4. Results

CSV Ingestion and Output Summary:

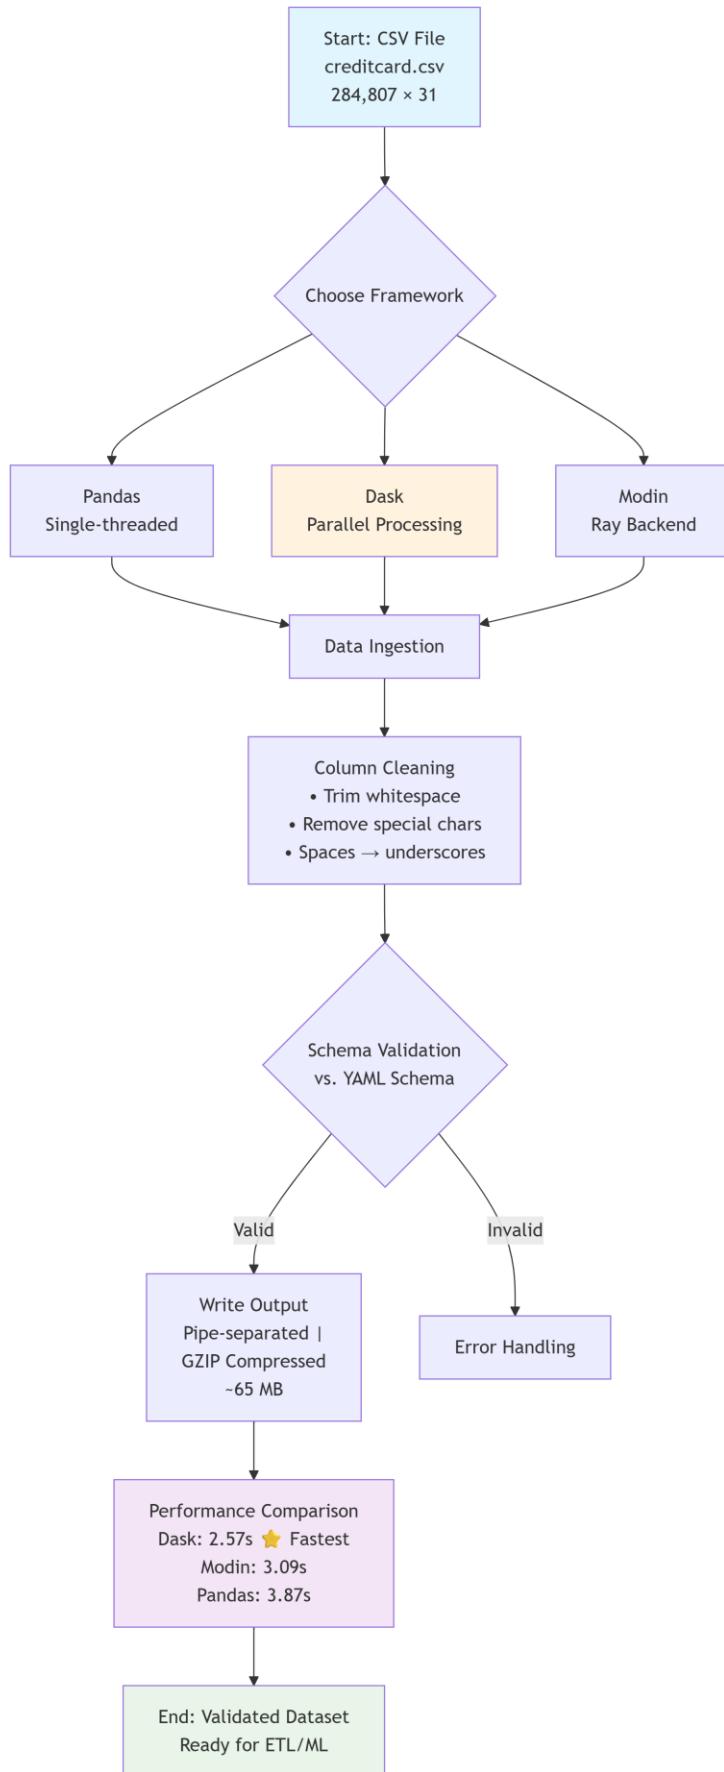
Framework	Rows	Columns	Time (s)	File Size (MB)
Pandas	284,807	31	3.87	65.41
Dask	284,807	31	2.57	65.48
Modin (Ray)	284,807	31	3.09	65.41

Observations:

- **Dask** was the fastest due to parallel computation.
- **Modin** provided an easy-to-use parallel Pandas alternative with similar results.
- All frameworks produced **validated and consistent outputs**.
- Compressed pipe-separated files are approximately **65 MB**, suitable for downstream processing.

5. Conclusion

- The assignment demonstrates **efficient file ingestion, cleaning, schema validation, and writing large datasets** using multiple frameworks.
- **Dask** was the most efficient for ingestion, while **all frameworks produced identical outputs**.
- The workflow can easily scale to **larger datasets**, making it practical for real-world ETL tasks.



Framework Performance Comparison

